

# Introduction to PIC Programming

## Baseline Architecture and Assembly Language

by David Meiklejohn, Gooligum Electronics

### Lesson 2: Flash an LED

In [lesson 1](#) we lit a single LED connected to one of the pins of a PIC10F200 or PIC12F508.

Now we'll make it flash.

In doing this, we will learn about:

- Using loops to create delays
- Variables
- Using exclusive-or (xor) to flip bits
- The ‘read-modify-write’ problem

The development environments and microcontrollers used for this lesson are the same as those in lesson 1.

Again, it is assumed that you are using a Microchip PICkit 2 or PICkit 3 programmer and either the [Gooligum Baseline and Mid-Range PIC Training and Development Board](#) or Microchip’s Low Pin Count (LPC) Demo Board, with Microchip’s MPLAB 8 or MPLAB X integrated development environment. But it is of course possible to adapt these instructions to a different programmers and/or development boards.

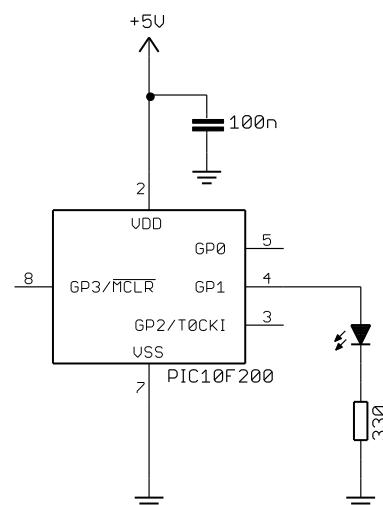
We will also assume that, if you have the Gooligum training board, you will continue to use the PIC10F200, and that if you have the Microchip LPC Demo Board, you will be using a PIC12F508 – both introduced in lesson 1.

### Example Circuit

Here's the PIC10F200 version of the circuit again.

If you have the Gooligum training board, simply plug the PIC10F200 into the 8-pin IC socket marked ‘10F’.

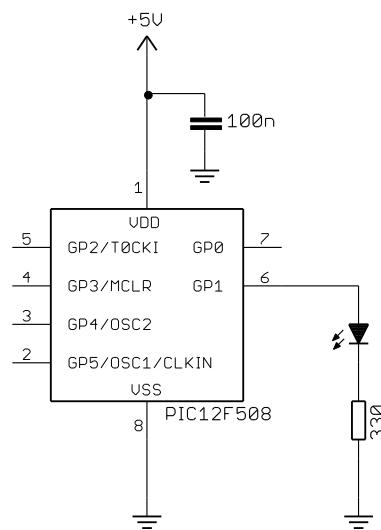
Connect a shunt across the jumper (JP12) on the LED labelled ‘GP1’, and ensure that every other jumper is disconnected.



Here's the corresponding PIC12F508 version.

You will need to use a PIC12F508 if you have Microchip's Low Pin Count Demo Board.

Refer back to [lesson 1](#) to see how to build this circuit, either by soldering a resistor, LED (and optional isolating jumper) to the demo board, or by making connections on the demo board's 14-pin header.



## Creating a new project

It is a good idea, where practical, to base a new software project on work you've done before. In this case, it makes sense to build on the program from lesson 1 – we just have to add extra instructions to flash the LED.

How to create a new project, based on an existing one, depends on whether you're using MPLAB 8 or MPLAB X, so we'll take a look at both.

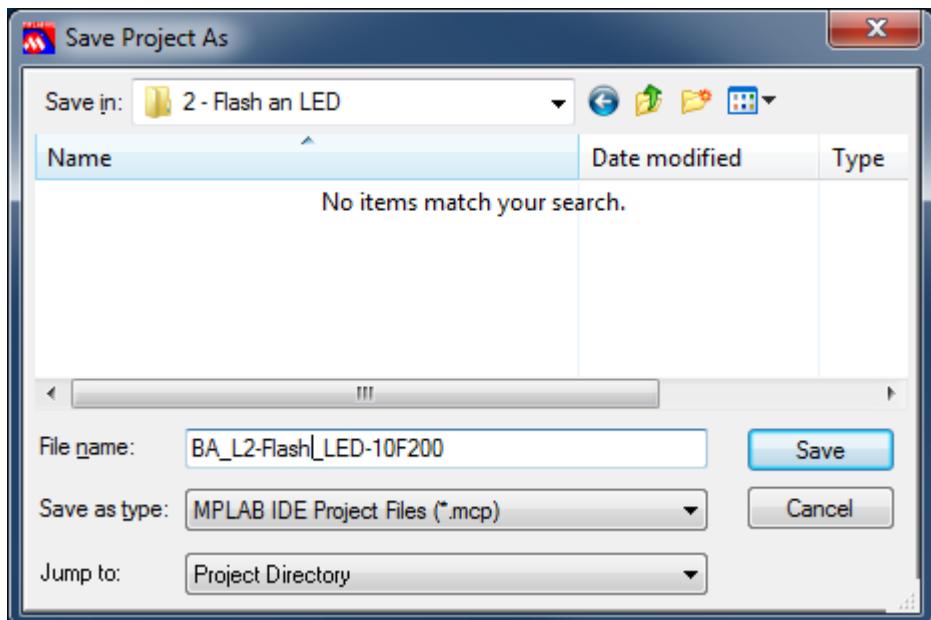
### MPLAB 8.xx

There are a couple of ways to do this, but the following method works well.

First, open the project you created in lesson 1 in MPLAB 8. You can do this easily by double-clicking the '\*.mcp' project file in your project folder.

Now use the "Project → Save Project As..." menu item to save the project in a new folder, with a new name.

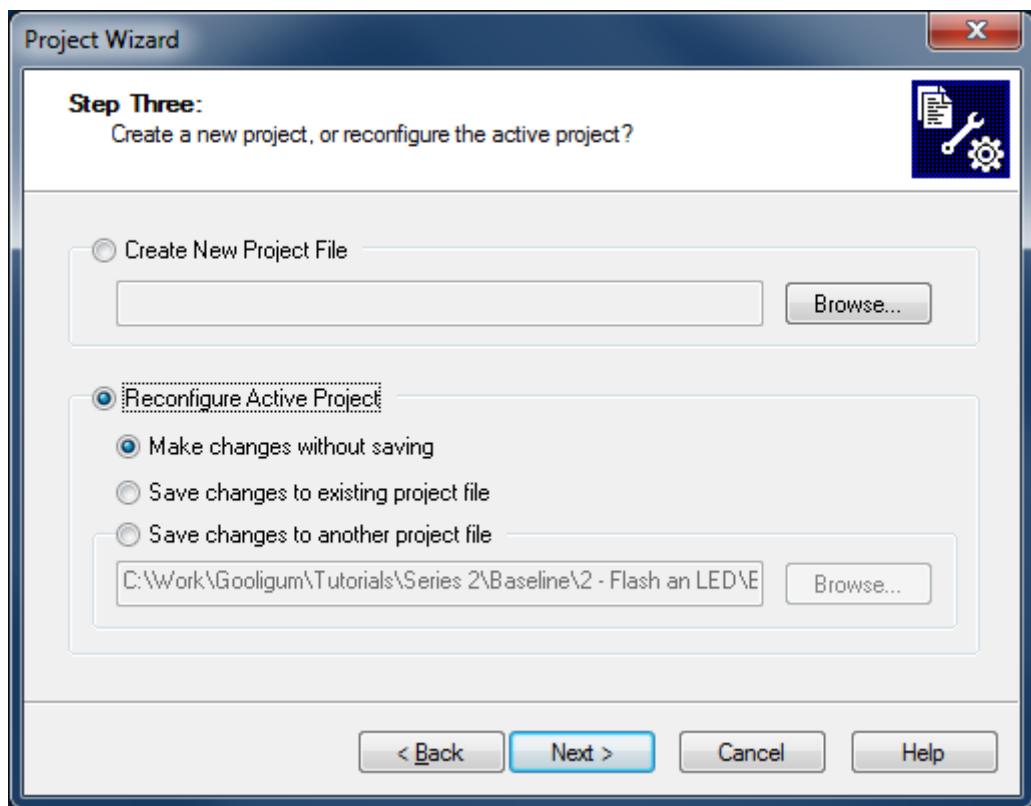
When a project is saved to a new location, all the files belonging to that project ("User" files, with relative paths) are copied to that location. You will find that in this case the '\*.asm' source file from lesson 1 has been copied into your new folder.



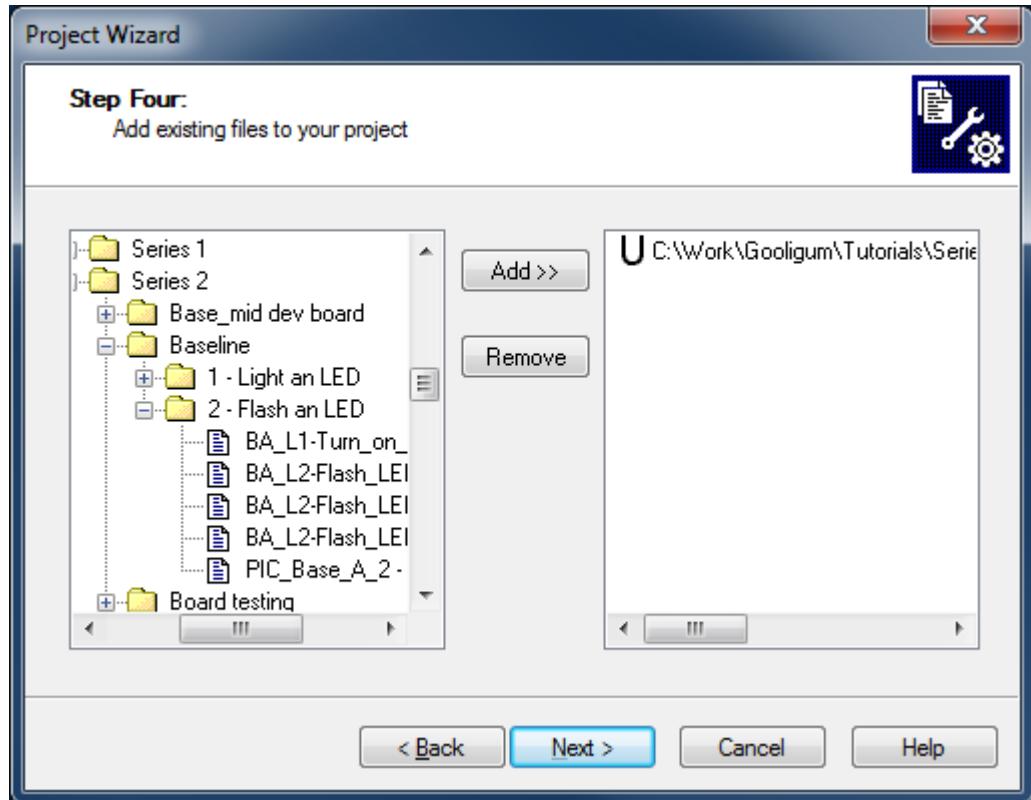
The next step is to use the project wizard (“Project → Project Wizard...”) to reconfigure the project, giving the source file a new name.

The correct device (PIC10F200 or PIC12F508) will already be selected, as will the toolsuite (MPASM), so simply click

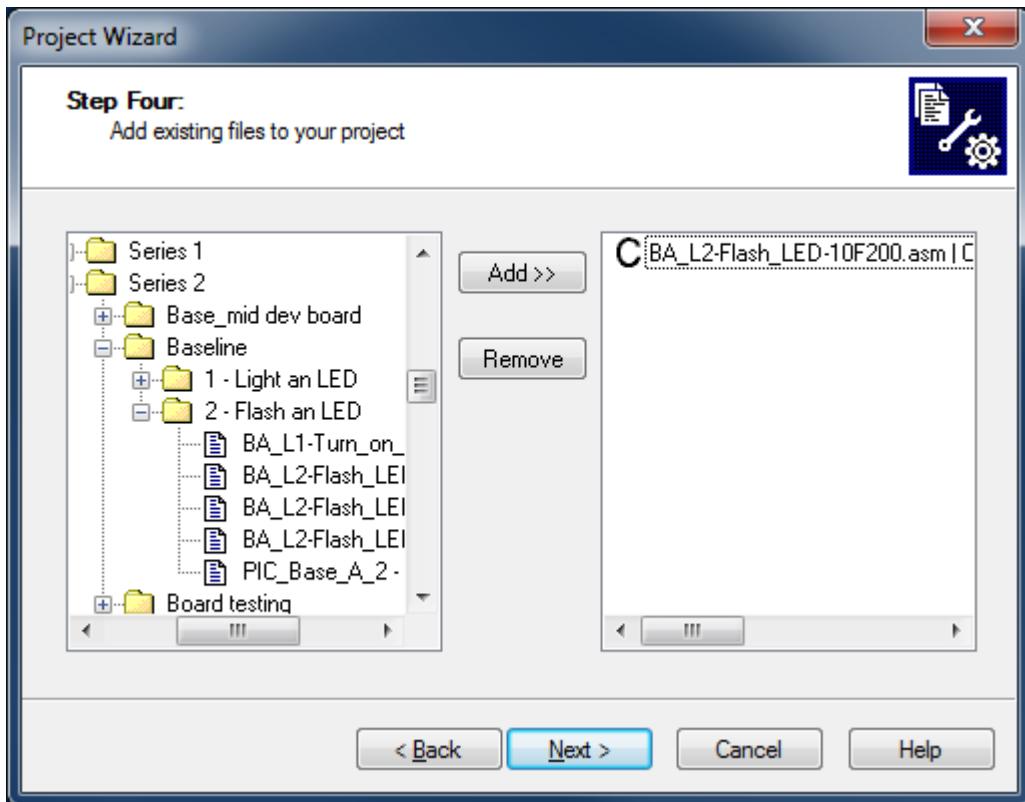
“Next” until you get to Step Three, and select “Reconfigure Active Project” and “Make changes without saving”, as shown:



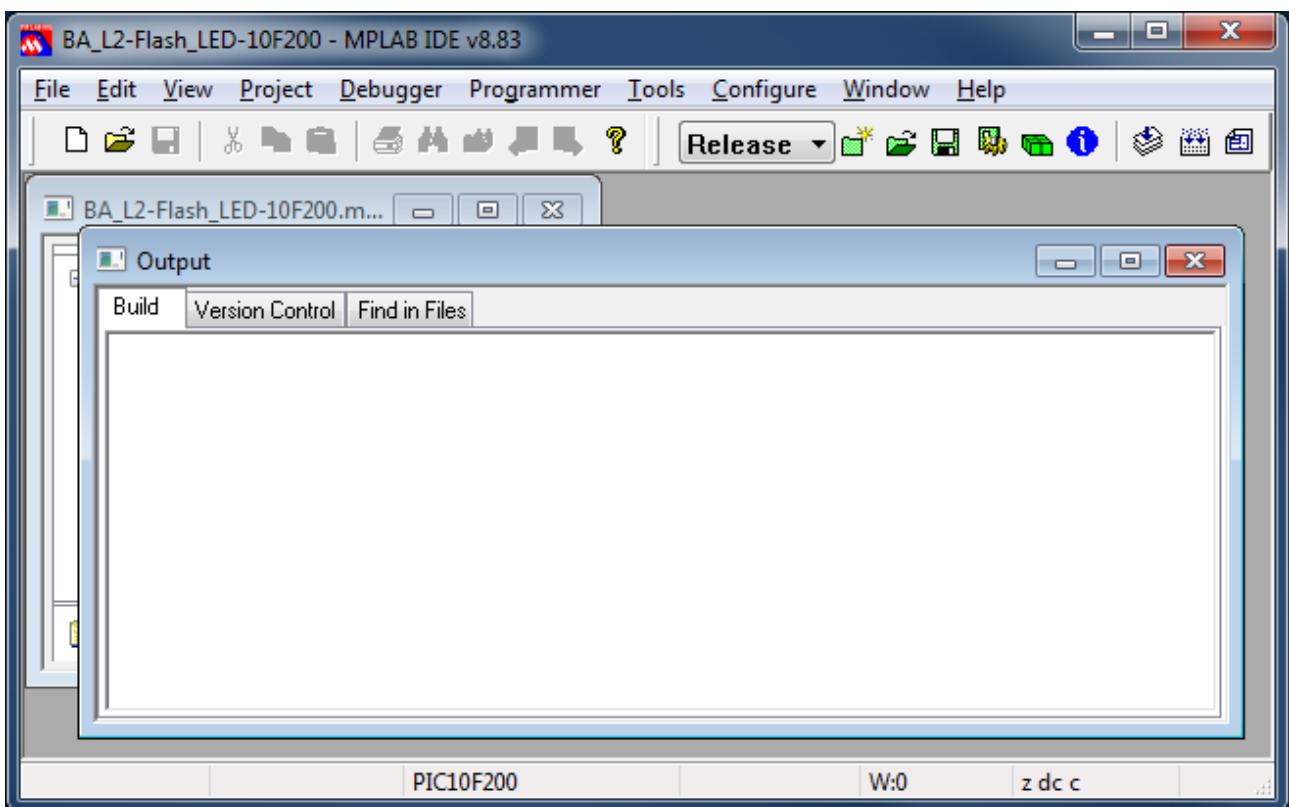
You are now presented with the following window, showing the assembler source file with a “U” to indicate a user file, in the new project directory, but with the same name as before:



Click on the “U” until it changes to a “C”. You can now click on the file name and rename it to something more appropriate to this lesson, such as ‘BA\_L2-Flash\_LED.asm’:

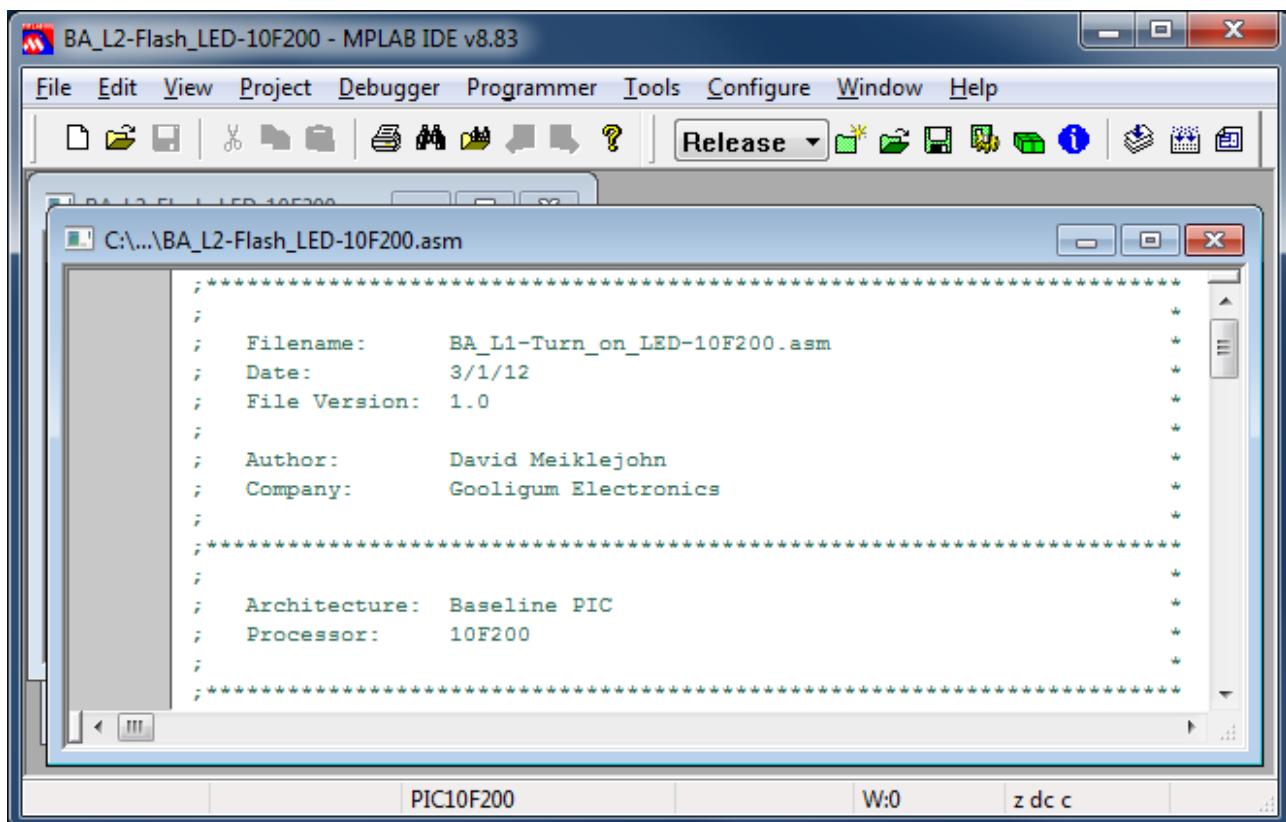


Finally click “Next” then “Finish” and the project is reconfigured, with the renamed source file:



It's a good idea at this point to save your new project, using the “Save Workspace” icon, or the “Project → Save Project” menu item.

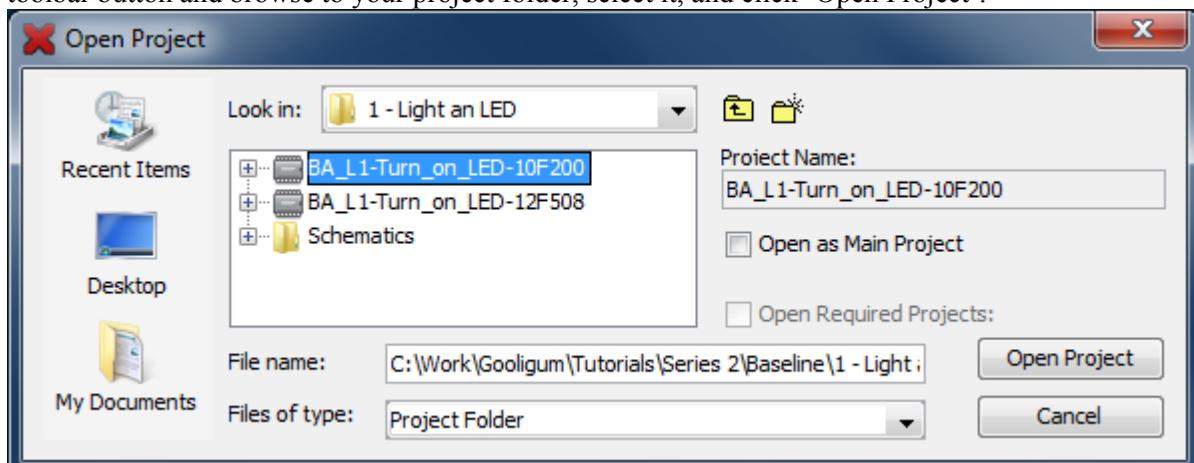
If you double-click on the source file (‘BA\_L2-Flash\_LED-10F200.asm’ in this example), you’ll see a copy of your code from lesson 1:



## MPLAB X

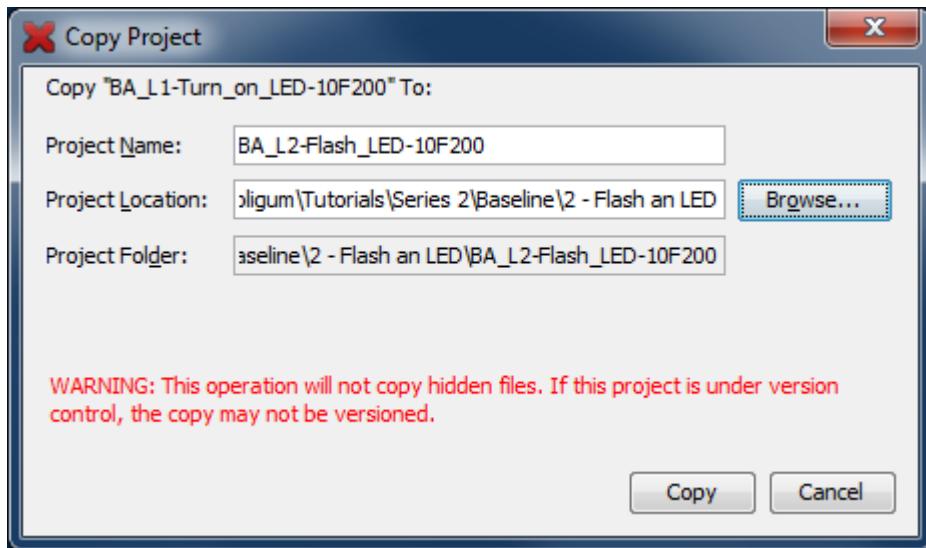
To create a new project in MPLAB X, based on an existing project, you first need to run MPLAB X (you can't simply double-click on a project file, like you can with MPLAB 8 projects), and then open your existing project within MPLAB X.

If you were recently working on the project you want to copy (such as the project from lesson 1), it is probably already visible in the Projects window. If it's not, it may appear under the “File → Open Recent Project” menu list. Or you can use the “File → Open Project” menu item, or click on the “Open Project...” toolbar button and browse to your project folder, select it, and click ‘Open Project’:



You should now right-click the project name ('BA\_L1-Turn\_on\_LED-10F200' in this example) in the Projects window, and select "Copy...".

The "Copy Project" dialog now gives you a chance to give your copied project a new name, such as 'BA\_L2-Flash\_LED'. You can also specify (and create, if you wish) a new folder for this project, by browsing to it:

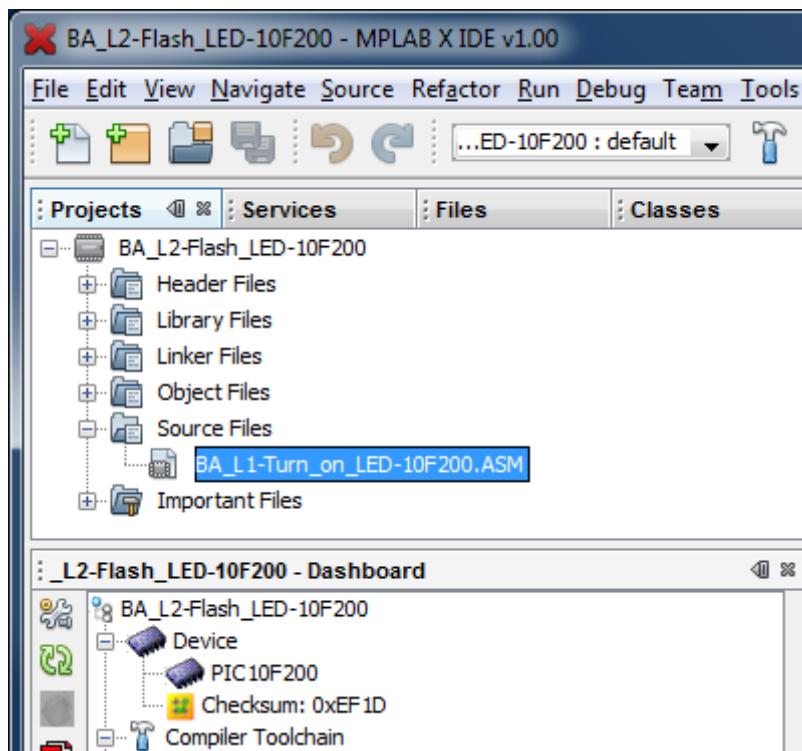


When you are satisfied with your new project name and location, click 'Copy'.

Your new project should now appear in the Projects window.

You can close your old project by right-clicking it and selecting "Close", so that only your new project is visible.

If you expand your new project, you'll see that source file from the old project has been copied into the new project, with its original name:



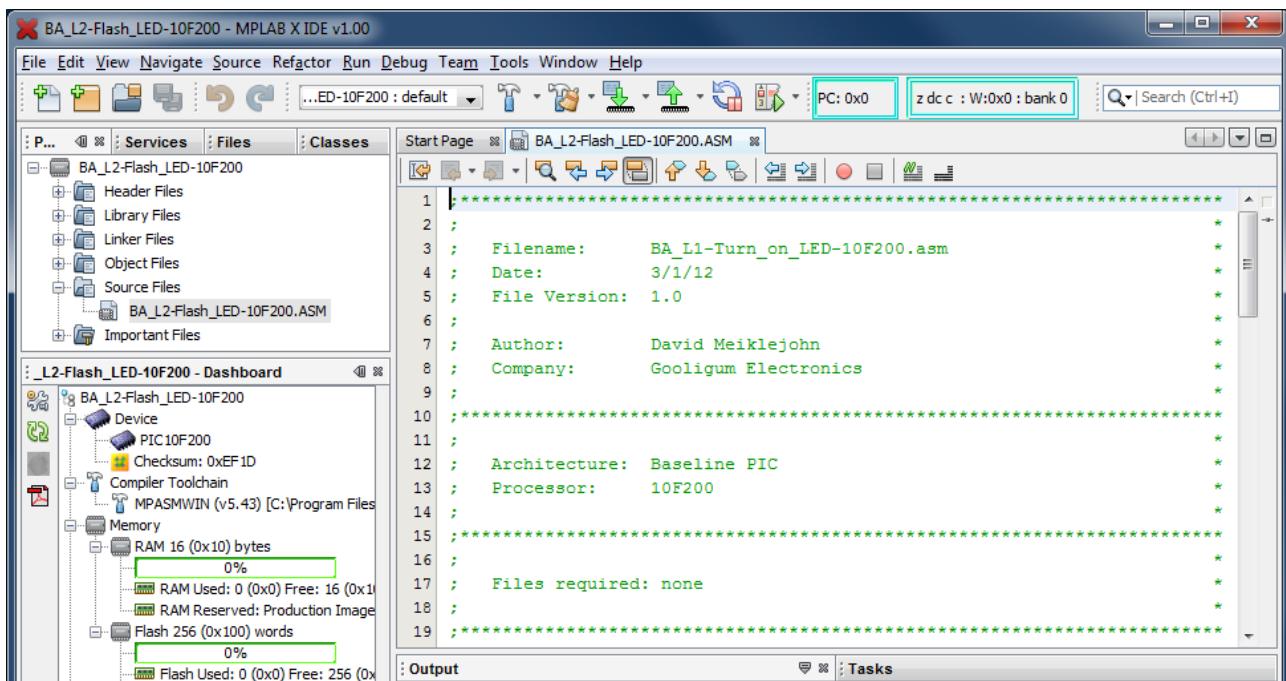
To rename the source file, to something more appropriate for this project, right-click it and select "Rename...".

Type in the new name, such as 'BA\_L2-Flash\_LED' and then click 'OK'.

Note that there is no need to type the '.ASM' suffix – the Rename dialog will keep the existing file extension.

You now have a new project, with a new name in a new location, with a renamed source file, copied from your old project.

If you double-click your new source file, you'll see a copy of your code from lesson 1 in an editor window:



## Flashing the LED

Whether you are using MPLAB 8 or X, you can now use the editor to update your code from lesson 1.

We'll need to add some code to make the LED flash, but first the comments should be updated to reflect the new project. For example:

```
;*****
;*
;* Filename:      BA_L2-Flash_LED-10F200.asm
;* Date:          20/1/12
;* File Version:  1.0
;*
;* Author:        David Meiklejohn
;* Company:       Gooligum Electronics
;*
;***** 
;*
;* Architecture: Baseline PIC
;* Processor:    10F200
;*
;***** 
;*
;* Files required: none
;*
;***** 
;*
;* Description:   Lesson 2, example 1
;*
;* Flashes a LED at approx 1 Hz.
;* LED continues to flash until power is removed.
;*
;*****
```

```
; Pin assignments:  
; GP1 = flashing LED  
;  
;*****
```

We're using the same PIC device as before, and it will be configured the same way, so we can leave the processor definition and configuration sections unchanged. There is also no need to change the internal RC oscillator calibration or reset vector sections.

So, for the PIC10F200 version, we still have, unchanged from lesson 1:

```
list      p=10F200  
#include <p10F200.inc>  
  
;***** CONFIGURATION  
; ext reset, no code protect, no watchdog  
_CONFIG _MCLRE_ON & _CP_OFF & _WDT_OFF  
  
;***** RC CALIBRATION  
RCCAL    CODE    0x0FF      ; processor reset vector  
        res 1          ; holds internal RC cal value, as a movlw k  
  
;***** RESET VECTOR *****  
RESET   CODE    0x000      ; effective reset vector  
        movwf OSCCAL    ; apply internal RC factory calibration
```

while for the PIC12F508, we have instead (also unchanged from lesson 1):

```
list      p=12F508  
#include <p12F508.inc>  
  
;***** CONFIGURATION  
; ext reset, no code protect, no watchdog, int RC clock  
_CONFIG _MCLRE_ON & _CP_OFF & _WDT_OFF & _IntRC_OSC  
  
;***** RC CALIBRATION  
RCCAL    CODE    0x1FF      ; processor reset vector  
        res 1          ; holds internal RC cal value, as a movlw k  
  
;***** RESET VECTOR *****  
RESET   CODE    0x000      ; effective reset vector  
        movwf OSCCAL    ; apply internal RC factory calibration
```

Again, we need to set up the PIC so that only GP1 is configured as an output, so we can leave the initialisation code from lesson 1 intact:

```
;***** MAIN PROGRAM *****  
  
;***** Initialisation  
start  
        movlw b'111101'      ; configure GP1 (only) as an output  
        tris    GPIO
```

In [Lesson 1](#), we made GP1 high, and left it that way. To make it flash, we need to set it high, then low, and then repeat. You may think that you could achieve this with something like:

```
flash
    movlw  b'000010'      ; set GP1 high
    movwf  GPIO
    movlw  b'000000'      ; set GP1 low
    movwf  GPIO
    goto   flash          ; repeat forever
```

If you try this code, you'll find that the LED appears to remain on continuously. In fact, it's flashing too fast for the eye to see.

Our PIC is using an internal RC oscillator<sup>1</sup>, clocked at a nominal 4 MHz. Each instruction executes in four clock cycles, or 1  $\mu$ s – except instructions which branch to another location, such as ‘goto’, which require two instruction cycles, or 2  $\mu$ s<sup>2</sup>.

This loop takes a total of 6  $\mu$ s, so the LED flashes at  $1/(6 \mu\text{s}) = 166.7 \text{ kHz}$ . That's much to fast to see!

To slow it down to a more sedate (and visible!) 1 Hz, we have to add a delay. But before looking at delays, we can make a small improvement to the code.

To flip, or toggle, a single bit – to change it from 0 to 1 or from 1 to 0, you can exclusive-or it with 1.

That is:

$$\begin{aligned} 0 \text{ XOR } 1 &= 1 \\ 1 \text{ XOR } 1 &= 0 \end{aligned}$$

So to repeatedly toggle GP1, we can read the current state of GPIO, exclusive-or the bit corresponding to GP1, then write it back to GPIO, as follows:

```
flash
    movlw  b'000010'      ; bit mask to toggle GP1 only
    xorwf  GPIO,f         ; toggle GP1 using mask in W
    goto   flash          ; repeat forever
```

The ‘xorwf’ instruction exclusive-ors the W register with the specified register – “exclusive-or **W** with file register”, and writes the result either to the specified file register (GPIO in this case) or to W.

Note that there is no need to set GP1 to an initial state; whether it's high or low to start with, it will be successively flipped.

Many of the PIC instructions, like xorwf, are able to place the result of an operation (e.g. add, subtract, or in this case XOR) into either a file register or W. This is referred to as the instruction destination. A ‘, f’ at the end indicates that the result should be written back to the file register; to place the result in W, use ‘, w’ instead.

This single instruction – ‘xorwf GPIO, f’ – is doing a lot of work. It reads GPIO, performs the XOR operation, and then writes the result back to GPIO.

### ***The read-modify-write problem***

And therein lays a potential problem. You'll find it referred to as the *read-modify-write* problem. When an instruction reads a port register, such as GPIO, the value that is read back is not necessarily the value that

<sup>1</sup> The 12F508 has been configured (using the \_\_config directive) to use its internal RC oscillator, while the 10F200 can only use an internal RC oscillator; there is no other choice.

<sup>2</sup> Assuming a 4 MHz processor clock

you originally wrote to it. When the PIC reads a port register, it doesn't read the value in the "output latch" (i.e. the value you wrote to it). Instead, it reads the pins themselves – the voltages present in the circuit.

Normally, that doesn't matter. When you write a '1', the corresponding pin (if configured as an output) will go to a high voltage level, and when you then read that pin, it's still at a high voltage, so it reads back as a '1'. But if there's excessive load on that pin, the PIC may not be able to drive it high, and it will read as a '0'. Or capacitance loading the output line may mean a delay between the PIC's attempt to raise the voltage and the voltage actually swinging high enough to register as a '1'. Or noise in the circuit may mean that a line that normally reads as a '1', sometimes (randomly) reads as a '0'.

In this simple case, particularly when we slow the flashing down to 1 Hz, you'll find that this isn't an issue. The above code will usually work correctly. But it's good to get into good habits early. For the reasons given above, it is considered "bad practice" to assume a value you have previously written is still present on an I/O port register.

It's better to keep a copy of what the port value is supposed to be, and operate on that, then copy it to the port register. This is referred to as using a *shadow register*.

We could use W as a shadow register, as follows:

```
    movlw  b'000000'      ; start with W zeroed
flash
    xorlw  b'000010'      ; toggle W bit corresponding to GP1 (bit 1)
    movwf  GPIO           ; and write to GPIO
    goto   flash          ; repeat forever
```

Each time around the loop, the contents of W are updated and then written to the I/O port.

The 'xorlw' instruction exclusive-ors a literal value with the W register, placing the result in W – "exclusive-or literal to W".

Normally, instead of 'movlw b'000000'' (or simply 'movlw 0') you'd use the 'clrw' instruction – "clear W".

'clrw' has the same effect as 'movlw 0', except that 'clrw' sets the 'Z' (zero) status flag, while the 'movlw' instruction doesn't affect any of the status flags, including Z.

Status flags are bits in the STATUS register:

	Bit 7	Bit 6	Bit 5	Bit 4	Bit 3	Bit 2	Bit 1	Bit 0
STATUS	GPWUF	-	-	TO	PD	Z	DC	C

Certain arithmetic or logical operations will set or clear the Z, DC or C status bits, and other instructions can test these bits, and take different actions depending on their value. We'll see examples of testing these flags in later lessons.

We're not using Z here, so we can use clrw to make the code more readable:

```
    clrw                  ; use W to shadow GPIO - initially zeroed
flash
    xorlw  b'000010'      ; toggle W bit corresponding to GP1 (bit 1)
    movwf  GPIO           ; and write to GPIO
    goto   flash          ; repeat forever
```

It would be very unusual to be able to use W as a shadow register, because it is used in so many PIC instructions. When we add delay code, it will certainly need to be able to change the contents of W, so we'll have to use a file register to hold the shadow copy of GPIO.

In [lesson 1](#), we saw how to allocate data memory for variables (such as shadow registers), using the UDATA and RES directives. In this case, we need something like:

```
;***** VARIABLE DEFINITIONS
UDATA
sGPIO    res 1                      ; shadow copy of GPIO
```

The flashing code now becomes:

```
clrf    sGPIO          ; clear shadow register
flash
  movf    sGPIO,w        ; get shadow copy of GPIO
  xorlw   b'000010'      ; toggle bit corresponding to GP1 (bit 1)
  movwf   sGPIO          ; in shadow register
  movwf   GPIO           ; and write to GPIO
  goto    flash          ; repeat forever
```

That's nearly twice as much code as the first version, that operated on GPIO directly, but this version is much more robust.

There are two new instructions here.

‘clrf’ clears (sets to 0) the specified register – “**clear file register**”.

‘movf’, with ‘, w’ as the destination, copies the contents of the specified register to W – “**move file register to destination**”. This is the instruction used to read a register.

‘movf’, with ‘, f’ as the destination, copies the contents of the specified register to itself. That would seem to be pointless; why copy a register back to itself? The answer is that the ‘movf’ instruction affects the Z status flag, so copying a register to itself is a sneaky way to test whether the value in the register is zero.

## Delay Loops

To make the flashing visible, we need to slow it down, and that means getting the PIC to “do nothing” between LED changes.

The baseline PICs do have a “do nothing” instruction: ‘nop’ – “**no operation**”. All it does is to take some time to execute.

How much time depends on the clock rate. Instructions are executed at one quarter the rate of the processor clock. In this case, the PIC is using the internal RC clock, running at a nominal 4 MHz (see [lesson 1](#)). The instructions are clocked at  $\frac{1}{4}$  of this rate: 1 MHz. Each instruction cycle is then 1  $\mu$ s.

Most baseline PIC instructions, including ‘nop’, execute in a single cycle. The exceptions are those which jump to another location (such as ‘goto’) or if an instruction is conditionally skipped (we’ll see an example of this soon). So ‘nop’ provides a 1  $\mu$ s delay – not very long!

Another “do nothing” instruction is ‘goto \$+1’. Since ‘\$’ stands for the current address, ‘\$+1’ is the address of the next instruction. Hence, ‘goto \$+1’ jumps to the following instruction – apparently useless behaviour. But all ‘goto’ instructions executes in two cycles. So ‘goto \$+1’ provides a 2  $\mu$ s delay in a single instruction – equivalent to two ‘nop’s, but using less program memory.

To flash at 1 Hz, the PIC should light the LED, wait for 0.5 s, turn off the LED, wait for another 0.5 s, and then repeat.

Our code changes the state of the LED once each time around the loop, so we need to add a delay of 0.5 s within the loop. That's 500,000  $\mu$ s, or 500,000 instruction cycles. Clearly we can't do that with 'nop's or 'goto's alone!

The answer, of course, is to use loops to execute instructions enough times to build up a useful delay. But we can't just use a 'goto', or else it would loop forever and the delay would never finish. So we have to loop some finite number of times, and for that we need to be able to count the number of times through the loop (incrementing or decrementing a loop counter variable) and test when the loop is complete.

Here's an example of a simple "do nothing" delay loop:

```
dly1    movlw   .10
        movwf   dc1           ; dc1 = 10 = number of loop iterations
        nop
        decfsz dc1,f
        goto   dly1
```

The first two instructions write the decimal value "10" to a loop counter variable called 'dc1'.

*Note that to specify a decimal value in MPASM, you prefix it with a '.'. If you don't include the '.', the assembler will use the default radix (hexadecimal), and you won't be using the number you think you are! Although it's possible to set the default radix to decimal, you'll run into problems if you rely on a particular default radix and then later copy and paste your code into another project, with a different default radix, giving different results. It's much safer, and clearer, to simply prefix all hexadecimal numbers with '0x' and all decimal numbers with '.'.*

The 'decfsz' instruction performs the work of implementing the loop – "decrement file register, skip if zero". First, it decrements the contents of the specified register, writes the result back to the register (as specified by the ', f' destination), then tests whether the result was zero. If it's not yet zero, the next instruction is executed, which will normally be a 'goto' which jumps back to the start of the loop. But if the result of the decrement is zero, the next instruction is skipped; since this is typically a 'goto', skipping it means exiting the loop.

The 'decfsz' instruction normally executes in a single cycle. But if the result is zero, and the next instruction is skipped, an extra cycle is added, making it a two-cycle instruction.

There is also an 'incfsz' instruction, which is equivalent to 'decfsz', except that it increments instead of decrementing. It's used if you want to count up instead of down. For a loop with a fixed number of iterations, counting down is more intuitive than counting up, so 'decfsz' is more commonly used for this.

In the code above, the loop counter, 'dc1', starts at 10. At the end of the first loop, it is decremented to 9, which is non-zero, so the 'goto' instruction is not skipped, and the loop repeats from the 'dly1' label. This process continues – 8,7,6,5,4,3,2 and on the 10<sup>th</sup> iteration through the loop, dc1 = 1. This time, dc1 is decremented to zero, and the "skip if zero" comes into play. The 'goto' is skipped, and execution continues after the loop.

You can see that the number of loop iterations is equal to the initial value of the loop counter (10 in this example). Call that initial number N. The loop executes N times.

To calculate the total time taken by the loop, add the execution time of each instruction in the loop:

nop	1
decfsz dc1,f	1 (except when result is zero)
goto dly1	2

That's a total of 4 cycles, except the last time through the loop, when the `decfsz` takes an extra cycle and the `goto` is not executed (saving 2 cycles), meaning the last loop iteration is 1 cycle shorter. And there are two instructions before the loop starts, adding 2 cycles.

Therefore the total delay time =  $(N \times 4 - 1 + 2)$  cycles =  $(N \times 4 + 1)$   $\mu$ s

If there was no ‘`nop`’, the delay would be  $(N \times 3 + 1)$   $\mu$ s; if two ‘`nop`’s, then it would be  $(N \times 5 + 1)$   $\mu$ s, etc.

It may seem that, because 255 is the highest 8-bit number, the maximum number of iterations ( $N$ ) should be 255. But not quite. If the loop counter is initially 0, then the first time through the loop, the ‘`decfsz`’ instruction will decrement it, and if an 8-bit counter is decremented from 0, the result is 255, which is non-zero, and the loop continues – another 255 times. Therefore the maximum number of iterations is in fact 256, with the loop counter initially 0.

So for the longest possible single loop delay, we can write something like:

```
dly1    clrf      dc1           ; loop 256 times
       nop
       decfsz   dc1,f
       goto     dly1
```

The two “move” instructions have been replaced with a single ‘`clrf`’, using 1 cycle less, so the total time taken is  $256 \times 4 = 1024$   $\mu$ s  $\approx 1$  ms.

That's still well short of the 0.5 s needed, so we need to wrap (or *nest*) this loop inside another, using separate counters for the inner and outer loops, as shown:

```
dly1    movlw    .200          ; loop (outer) 200 times
       movwf    dc2
       clrf     dc1           ; loop (inner) 256 times
       nop      ; inner loop = 256 x 4 - 1 = 1023 cycles
       decfsz  dc1,f
       goto    dly1
       decfsz  dc2,f
       goto    dly1
```

The loop counter ‘`dc2`’ is being used to control how many times the inner loop is executed.

Note that there is no need to clear the inner loop counter (`dc1`) on each iteration of the outer loop, because every time the inner loop completes, `dc1 = 0`.

The total time taken for each iteration of the outer loop is 1023 cycles for the inner loop, plus 1 cycle for the ‘`decfsz dc2,f`’ and 2 cycles for the ‘`goto`’ at the end, except for the final iteration, which, as we've seen, takes 1 cycle less. The three setup instructions at the start add 3 cycles, so if the number of outer loop iterations is  $N$ :

Total delay time =  $(N \times (1023 + 3) - 1 + 3)$  cycles =  $(N \times 1026 + 2)$   $\mu$ s.

The maximum delay would be with  $N = 256$ , giving 262,658  $\mu$ s. We need a bit less than double that. We could duplicate all the delay code, but it takes fewer lines of code if we duplicate only the inner loop:

```
; delay 500ms
dly1    movlw    .244          ; outer loop: 244 x (1023 + 1023 + 3) + 2
       movwf    dc2
       clrf     dc1           ; inner loop: 256 x 4 - 1
       nop      ; inner loop 1 = 1023 cycles
       decfsz  dc1,f
       goto    dly1
dly2    nop      ; inner loop 2 = 1023 cycles
       decfsz  dc1,f
       goto    dly2
       decfsz  dc2,f
       goto    dly1
```

The two inner loops of 1023 cycles each, plus the 3 cycles for the outer loop control instructions (`decfsz` and `goto`) make a total of 2049 µs. Dividing this into 500,000 gives 244.02 – pretty close to a whole number, so an outer loop count of 244 will be very close to what's needed.

The calculations are shown in the comments above. The total time for this delay code is 499,958 cycles. In theory, that's 499.958 ms – within 0.01% of the desired result! Given that that's much more accurate than the 4 MHz internal RC oscillator, there is no point trying for more accuracy than this.

But suppose the calculation above had come out as needing some fractional number of outer loop iterations, say 243.5 – what would you do? Generally you'd fine-tune the timing by adding or removing 'nop's. E.g. suppose that both inner loops had 2 'nop's instead of 1. Then they would execute in  $256 \times 5 - 1 = 1279$  cycles, and the calculation for the outer loop counter would be  $500,000 \div (1279 + 1279 + 3) = 195.24$ . That's not as good a result as the one above, because ideally we want a whole number of loops. 244.02 is much closer to being a whole number than 195.24.

For even finer control, you can add 'nop's to the outer loop, immediately before the '`decfsz dc2, f`' instruction. One extra 'nop' would give the outer loop a total of  $1023 + 1023 + 4 = 2050$  cycles, instead of 2049. The loop counter calculation becomes  $500,000 \div 2050 = 243.90$ . That's not bad, but 244.02 is better, so we'll leave the code above unchanged.

With a bit of fiddling, once you get some nested loops close to the delay you need, adding or removing 'nop' or '`goto $+1`' instructions can generally get you quite close to the delay you need. And remember that it is pointless to aim for high precision (< 1%) when using the internal RC oscillator. When using a crystal, it makes more sense to count every last cycle accurately, as we'll see in [lesson 7](#).

For delays longer than about 0.5 s, you'll need to add more levels of nesting to your delay loops – with enough levels you can count for years!

### ***Complete program***

Putting together all these pieces, here's the complete PIC10F200 version of our LED flashing program:

```
;***** CONFIGURATION
;
; Description:    Lesson 2, example 1
;
; Flashes a LED at approx 1 Hz.
; LED continues to flash until power is removed.
;
;***** VARIABLE DEFINITIONS
UDATA
sGPIO    res 1           ; shadow copy of GPIO
```

```

dc1      res 1          ; delay loop counters
dc2      res 1

;***** RC CALIBRATION
RCCAL    CODE    0x0FF      ; processor reset vector
          res 1          ; holds internal RC cal value, as a movlw k

;***** RESET VECTOR ****
RESET   CODE    0x000      ; effective reset vector
          movwf   OSCCAL     ; apply internal RC factory calibration

;***** MAIN PROGRAM ****
;***** Initialisation
start
          movlw   b'111101'    ; configure GP1 (only) as an output
          tris    GPIO

          clrf    sGPIO        ; start with shadow GPIO zeroed

;***** Main loop
main_loop
          ; toggle LED on GP1
          movf    sGPIO,w      ; get shadow copy of GPIO
          xorlw   b'000010'      ; toggle bit corresponding to GP1 (bit 1)
          movwf   sGPIO         ; in shadow register
          movwf   GPIO          ; and write to GPIO

          ; delay 500ms
          movlw   .244          ; outer loop: 244 x (1023 + 1023 + 3) + 2
          movwf   dc2            ; = 499,958 cycles
          clrf    dc1            ; inner loop: 256 x 4 - 1
dly1    nop              ; inner loop 1 = 1023 cycles
          decfsz dc1,f
          goto   dly1

dly2    nop              ; inner loop 2 = 1023 cycles
          decfsz dc1,f
          goto   dly2
          decfsz dc2,f
          goto   dly1

          goto   main_loop      ; repeat forever

END

```

The 12F508 version is very similar, with changes to the `list`, `#include`, `__CONFIG` and `RCCAL CODE` directives, as shown earlier.

If you follow the programming procedure described in [lesson 1](#), you should now see your LED flashing at something very close to 1 Hz.

## Conclusion

It's taken two lessons and dozens of pages to get here, but we finally have a flashing LED!

In this lesson, we built on the first, showing how to base a new project on an existing one, modifying it and adding whatever additional features the new project needs.

We saw how to toggle a pin, discussed how “read-modify-write” operations on a port can be problematic, and showed how to use shadow registers can be used to avoid such potential problems.

We also saw how to use decrement instructions with conditional tests to implement loops, and how to use loops to create delays of any length.

In the [next lesson](#) we'll step up to a slightly bigger PIC, the 12F509.

We'll also see how to make our programs more modular, so that useful pieces of code such as the 500 ms delay developed here can be easily re-used.